



TEXAS ADVANCED COMPUTING CENTER

WWW.TACC.UTEXAS.EDU



TEXAS

The University of Texas at Austin

Manycore Optimization: Vectorization

Todd Evans

<http://tinyurl.com/tacc-manycore-2017>

February 16, 2018

3 Rules of Thumb for Optimization

Applies to Serial, Vectorized, and Multithreaded Code

1. Make compute intensive parts of code clean
 - 1.1 Compiler likes linear, predictable data structures
 - 1.2 Save OO for staging and managing workflow
 - 1.3 Compiler will do most of the work for you
2. Use optimization options (flags)
 - 2.1 these tell compiler what optimization are allowed
 - 2.2 these tell compiler what instructions are allowed
3. Use math libraries (lapack/blas, MKL)
 - 3.1 These may thread automagically for you
 - 3.2 Teams of engineers have developed these over decades
 - 3.3 Specialized for particular architectures (e.g. KNL)

How-to Vectorize Your Program

Approaches

- ▶ Assembly (Bad)
 - ▶ Error prone, probably won't be that fast, not portable
- ▶ Intrinsics (Better)
 - ▶ Less error prone, probably not the fastest possible, still not portable
- ▶ Compiler (Good)
 - ▶ This is often the best approach!
 - ▶ Portable
 - ▶ Automatically vectorizes your code
- ▶ Link to optimized, vectorized libraries if available (Very good!)
 - ▶ Teams of scientists and engineers have put years into optimizing these
 - ▶ Portable
 - ▶ MKL

Vector Aware Programming

- ▶ Most vectorization opportunities are in loops
- ▶ Keep loops simple to help compiler
- ▶ Evaluate compiler output (tools exist for this)
- ▶ Evaluate Performance Improvement
- ▶ Consider Memory Access Patterns (unit stride is good)
stride is the number of memory locations between accessed elements of an array
- ▶ Vector Math Libraries should be used for special functions
- ▶ MKL or equivalent should be used for linear algebra

How-to Achieve Automatic Loop Vectorization

Compiler flags - Loops will be vectorized by compiler

- ▶ Skylake and KNL support AVX512 but there are 9 subsets - they support 2 in common
- ▶ Intel
 - ▶ -O2 (SSE2/3/4)
 - ▶ -xhost (chooses most advanced)
 - ▶ -xavx/-xavx2/-xmic-avx512/-xcore-avx512
- ▶ GCC
 - ▶ -O3 (SSE2/3/4)
 - ▶ -march=native (chooses most advanced)
 - ▶ -mavx/-mavx2/-mavx512f -mavx512cd -mavx512er
-mavx512pf/-mavx512f -mavx512cd -mavx512bw -mavx512dq
-mavx512vl -mavx512ifma -mavx512vbmi

How-to Achieve Loop Vectorization

Characteristics required for loop vectorization

- ▶ Countable
 - ▶ Number of iterations must be known before loop executes
 - ▶ No conditional termination
- ▶ Single flow control
 - ▶ No switch statements
 - ▶ if statements are possible if masked
- ▶ Only innermost loop is vectorized
- ▶ No function calls allowed (except inlined or VML)
- ▶ No true loop *dependencies*

Loop Dependencies

- ▶ Vectorization changes order of computation from sequential
- ▶ Compiler wants to vectorize but is careful to ensure correct results (conservative)
- ▶ Compiler performs dependency analysis
 - ▶ Dependent on vector width
- ▶ Compiler performs cost analysis
 - ▶ Not always faster to vectorize
 - ▶ Overhead associated with checking dependencies at runtime
 - ▶ short loops not always worth it

Dependency Analysis

Read After Write (RAW)

Variable is written then read

Not Vectorizable

Intel calls this FLOW

```
for (i=1; i<N; i++)  
    a[i]=a[i-1]+b[i];
```

Write After Read (WAR)

Variable is read then written

Vectorizable

Intel calls this ANTI

```
for (i=1; i<N-1; i++)  
    a[i]=a[i+1]+b[i];
```


Dependency Analysis

Read After READ (RAR)

Variable is read then read

Vectorizable

Not really a dependency

```
for (i=1; i<N; i++)  
    a[i]=b[i%M]+c[i];
```

Write After Write (WAW)

Variable is written then written

NOT Vectorizable

```
for (i=1; i<N-1; i++)  
    a[i%M]=b[i]+c[i];
```

Read After Write

Dependency Analysis: Do loads, operations, and stores across iterations in parallel produce the same results as executing each iteration sequentially?

```
int a[5] = {1, 1, 1, 1, 1}
int b[5] = {2, 2, 2, 2, 2}

for (i=1; i<N; i++)
    a[i]=a[i-1]*b[i];
```

iteration 1	a(2)	=	a(1)	*	b(2)
iteration 2	a(3)	=	a(2)	*	b(3)
iteration 3	a(4)	=	a(3)	*	b(4)
iteration 4	a(5)	=	a(4)	*	b(5)
Store a(2-5)		Load a(1-4)	operate	Load b(2-5)	

Read After Write

Dependency Analysis: Do loads, operations, and stores across iterations in parallel produce the same results as executing each iteration sequentially? NO!

Sequential

iteration 1	$a(2)=2$	=	$a(1)=1$	*	$b(2)=2$
iteration 2	$a(3)=4$	=	$a(2)=2$	*	$b(3)=2$
iteration 3	$a(4)=8$	=	$a(3)=4$	*	$b(4)=2$
iteration 4	$a(5)=16$	=	$a(4)=8$	*	$b(5)=2$

Vectorized

iteration 1	$a(2)=2$	=	$a(1)=1$	*	$b(2)=2$
iteration 2	$a(3)=2$	=	$a(2)=1$	*	$b(3)=2$
iteration 3	$a(4)=2$	=	$a(3)=1$	*	$b(4)=2$
iteration 4	$a(5)=2$	=	$a(4)=1$	*	$b(5)=2$

Write After Read

Dependency Analysis: Do loads, operations, and stores across iterations in parallel produce the same results as executing each iteration sequentially?

```
for (i=1; i<N-1; i++)  
    a[i]=a[i+1]*b[i];
```

iteration 1	a(1)	=	a(2)	*	b(1)
iteration 2	a(2)	=	a(3)	*	b(2)
iteration 3	a(3)	=	a(4)	*	b(3)
iteration 4	a(4)	=	a(5)	*	b(4)
Store a(1-4)		Load a(2-5)		operate	Load b(1-4)

Write After Read

Dependency Analysis: Do loads, operations, and stores across iterations in parallel produce the same results as executing each iteration sequentially? Yes!

Sequential

iteration 1	$a(1)=2$	=	$a(2)=1$	*	$b(1)=2$
iteration 2	$a(2)=2$	=	$a(3)=1$	*	$b(2)=2$
iteration 3	$a(3)=2$	=	$a(4)=1$	*	$b(3)=2$
iteration 4	$a(4)=2$	=	$a(5)=1$	*	$b(4)=2$

Vectorized

iteration 1	$a(1)=2$	=	$a(2)=1$	*	$b(1)=2$
iteration 2	$a(2)=2$	=	$a(3)=1$	*	$b(2)=2$
iteration 3	$a(3)=2$	=	$a(4)=1$	*	$b(3)=2$
iteration 4	$a(4)=2$	=	$a(5)=1$	*	$b(4)=2$

Striding: Cache Reuse


Distance between elements accessed in memory

- ▶ unit stride: `for(i=0;i<N;i++) a[i]=b[i];`
stride of 2: `for(i=0;i<N;i+=2) a[i]=b[i];`
stride of 4: `for(i=0;i<N;i+=4) a[i]=b[i];`
- ▶ whole cache lines (8 DP elements) are always moved from memory to cache
 - ▶ slow to move data from memory to cache
 - ▶ unit stride uses each element (enhances cache reuse)
 - ▶ unit stride allows 8 DP to be loaded from cache into vector register w/ 1 load!
 - ▶ non-unit stride requires multiple loads or gather/scatter :(
- ▶ effect is greatest when reading from memory, least when reading from L1

Why use caches?

Memory Hierarchy

- ▶ Caches are smaller but faster than memory
- ▶ May have multiple levels of cache: L1, L2 ...
- ▶ Helps keep cores fed with data
- ▶ Cache reuse is often critical to performance



Striding

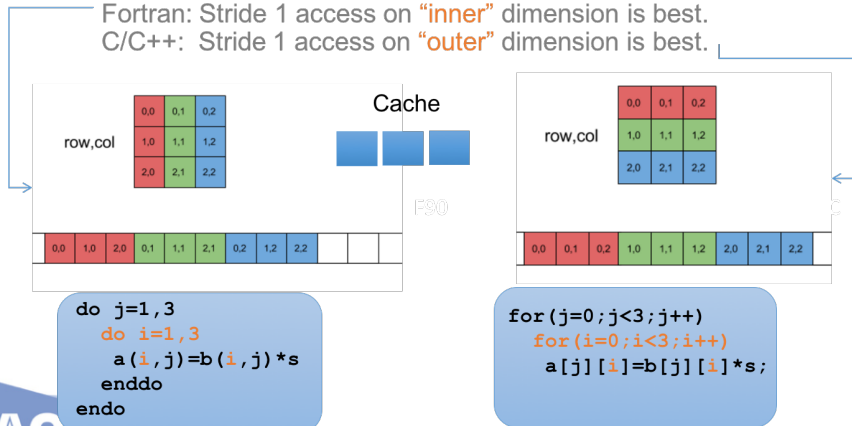
1-D Array

Stride 1 access is best! Uses all the elements in the cache line!

Multi-D Array

Fortran: Stride 1 access on “inner” dimension is best.

C/C++: Stride 1 access on “outer” dimension is best.




Strided Access

```
*do i = 1,4000000*istride,istride  
  a(i) = b(i) + c(i) * sfactor  
enddo
```

Memory BW is scarce in Manycore


- ▶ Striding is even more important
- ▶ Striding through memory is bad
 - ▶ stride 2 = 1/2 effective bandwidth
 - ▶ stride 4 = 1/4 effective bandwidth
 - ▶ stride 8 = 1/8 effective bandwidth
- ▶ Striding through L2 and L1 not *as* bad
 - ▶ still requiring more loads

Memory Strided Add* Performance



```
*do i = 1,2048*istride,istride  
  a(i) = b(i) + c(i) * sfactor  
enddo
```



L2 Strided Add* Performance



Alignment

- ▶ Data objects can be created in memory on specific byte boundaries
- ▶ Can increase efficiency of loads and stores
 - ▶ A vector load can only access one cache line at a time (gather accesses multiple but over several cycles)
- ▶ For AVX/AVX2 (SNB/HSW), alignment to 32B boundaries
 - ▶ allows a single load to move 4 DP data elements
- ▶ For AVX-512/MIC-AVX-512 (KNL), alignment to 64B boundaries
 - ▶ allows a single load to move 8 DP data elements
- ▶ Compiler can detect lack of alignment and peel off iterations to achieve optimal alignment for bulk of loop
 - ▶ Some overhead
 - ▶ `void *memalign(size_t alignment, size_t size);`

Alignment: Loads



Alignment: Padding

In multi-dimensional array, resizing (padding) lower dimensions for alignment can improve performance

$a[4][3] \rightarrow a[4][4]$



```
for(i=0; i<nx; i++)  
  for(j=0; j<ny; j++)  
    b[i][j] = a[i][j];
```

Inlining

Functions within loops prevent vectorization

- ▶ Inlining can overcome this
- ▶ If function call and def are in same file, `-ip` or `-O3 inlines`
- ▶ If function call and def are in separate file, only `-ipo` inlines
- ▶ After inlining vectorization occurs

```
for(i=0;i<nx;i++) {  
    x=x0+i*h;  
    sum=sum+  
        do_r2(x,y,xp,yp);  
}
```

```
double do_r2(double x,double y,  
             double xp, double yp) {  
    double r2;  
    r2=(x-xp)*(x-xp)+  
        (y-yp)*(y-yp);  
    return r2;  
}
```

Inlining effects

Before Inlining

```
double r2(x,x0,y,y0)
    return (x-x0)*(x-x0)+(y-y0)*(y-y0);

int main() {
    ...
    for (i=0;i<Nx;i++)
        d[i]+=r2(x[i],x0[i],y[i],y0[i]);
    ...
    return 0;
}
```

After Inlining

```
int main() {
    ...
    for (i=0;i<Nx;i++)
        d[i]+=(x[i]-x0[i])*(x[i]-x0[i])+\
            (y[i]-y0[i])*(y[i]-y0[i]));
    ...
    return 0;
}
```

Inlining	Vectorization	Time (ms)
no	no	1.55
yes	no	0.44
yes	yes	0.056

Aliasing

- ▶ Two different pointers can have the same target (e.g. point to the same array or array element)
- ▶ Loops in routines that use pointer references might not vectorize
- ▶ Compiler has to be certain vectorization does not introduce dependencies
- ▶ In functions that pass pointers, ensure the compiler can detect whether references overlap

Example: `my_cp(N, &b[1], &b[0])` would be a RAW

```
void my_cp(int N, double *a, double *b) {  
    for (int i=0; i<N-1; i++)  
        a[i]=b[i];  
}
```

Aliasing

Compiler can test for overlap in simple circumstances

```
void my_cp(int N, double *a, double *b) {  
    if (a+N>b || b+N>a)  
        for (int i=0; i<N; i++)  
            a[i]=b[i];  
    else  
        for (i=0; i<N; i+=4){  
            a[i+0]=b[i+0];  
            a[i+1]=b[i+1];  
            a[i+2]=b[i+2];  
            a[i+3]=b[i+3];  
        }  
}
```


Aliasing

Overlap test can be too expensive in some cases

```
void my_cp(int nx, double *a, double *b, int *ioff) {  
    for (int i=0; i<nx; i++)  
        a[i]=b[i+*ioff];  
}
```

- ▶ Compiler decides runtime overlap test is not cost effective
- ▶ `b` may be aliased to `a` and/or `*ioff` might point to `a[i]`
- ▶ multiple possible dependencies
- ▶ `-ansi-alias` flag prevents compiler from aliasing different types (enables vectorization in this example)

Pragmas

Pragmas can also tell compiler no aliasing: Developer's responsibility

- ▶ `#pragma ivdep`
 - ▶ no assumed dependencies
 - ▶ vectorize if dependency not proven
- ▶ `#pragma vector always`
 - ▶ vectorize if no assumed or proven dependencies
 - ▶ ignore cost
- ▶ `#pragma novector`
 - ▶ never vectorize
- ▶ `#pragma simd`
 - ▶ always vectorize
 - ▶ ignore cost and safety

Vectorization Summary

- ▶ KNL with MIC-AVX512 is here: $(512\text{b width} \times \text{FMA}) = 16$ simultaneous DP FLOPs per cycle (and 2 VPU's so 32!)
- ▶ SKX with CORE-AVX512 is too!
- ▶ Wider vector register does not guarantee improved performance
- ▶ Ensure no explicit dependencies
- ▶ Ensure no functions (or inline)
- ▶ Ensure no aliasing

Compiler Generated Optimization Reports

The Intel compiler can generate optimization reports

- ▶ vectorization reports
 - ▶ what loops were or were not vectorized and why
 - ▶ vectorization count
 - ▶ peeled and remainder loops
- ▶ OpenMP parallelization
- ▶ inlining
- ▶ loop transformations
 - ▶ vectorization
 - ▶ cache reuse

Compiler Generated Optimization Reports

Pretty easy to use

- ▶ `$ icc -xhost -qopt-report=5 example.c -o example`
- ▶ Two main sections:

- ▶ Report from: Interprocedural optimizations [ipo]

```
→ EXTERN: (23,23) _vla_alloc(long)
→ EXTERN: (34,7) clock(void)
→ INLINE: (35,12) rdtsc() (isz = 4) (sz = 9)
→ INLINE: (43,12) rdtsc() (isz = 4) (sz = 9)
```

- ▶ Report from: Loop nest, Vector & Auto-parallelization optimizations [loop, vec, par]

```
LOOP BEGIN at vector.c(39,5)
remark #25444: Loopnest Interchanged: ( 1 2 3 ) --> ( 2 1 3 )
remark #15542: loop was not vectorized: inner loop was already vectorized [ vector.c(39,\
5) ]
```

Compiler Generated Optimization Reports

```
LOOP BEGIN at omp_vector.c(46,7)
  <Peeled loop for vectorization>
    remark #15389: vectorization support: reference B[i][j] has unaligned access [ omp_vector.c(47,10) ]
    remark #15381: vectorization support: unaligned access used inside loop body
    remark #15335: peel loop was not vectorized: vectorization possible but seems inefficient. Use
      vector always directive or -vec-threshold0 to override
    remark #15305: vectorization support: vector length 8
    remark #15309: vectorization support: normalized vectorization overhead 2.158
    remark #25015: Estimate of max trip count of loop=7
  LOOP END
```

```
LOOP BEGIN at omp_vector.c(46,7)
  remark #25085: Preprocess Loopnests: Moving Out Load and Store [ omp_vector.c(47,2) ]
  remark #15388: vectorization support: reference B[i][j] has aligned access [ omp_vector.c(47,10) ]
  remark #15305: vectorization support: vector length 8
  remark #15399: vectorization support: unroll factor set to 8
  remark #15309: vectorization support: normalized vectorization overhead 0.643
  remark #15300: LOOP WAS VECTORIZED
  remark #15448: unmasked aligned unit stride loads: 1
  remark #15475: — begin vector cost summary —
  remark #15476: scalar cost: 7
  remark #15477: vector cost: 0.870
  remark #15478: estimated potential speedup: 6.100
  remark #15488: — end vector cost summary —
LOOP END
```

```
LOOP BEGIN at omp_vector.c(46,7)
  <Remainder loop for vectorization>
    remark #15388: vectorization support: reference B[i][j] has aligned access [ omp_vector.c(47,10) ]
    remark #15305: vectorization support: vector length 8
    remark #15309: vectorization support: normalized vectorization overhead 2.111
```

Additional Considerations: Skylake

Skylake processors use dynamic frequency scaling

- ▶ Different Width Instructions run at different frequencies
 - ▶ SSE (128-bit) frequency is higher than AVX2 (256-bit) is higher than AVX512 (512-bit)
 - ▶ Different is significant: 2.0 - 2.8 when using all cores
- ▶ Intel compiler generates instructions that use lower half of 512-bit registers preferentially
 - ▶ CORE-AVX512 instructions have extensions that are 128-bit and 256-bit in addition to the nominal 512-bit width
- ▶ `-qopt-zmm-usage=high` tells compiler to more aggressively use full 512-bit
 - ▶ good for code that vectorizes well
 - ▶ available with Intel 18 compiler

License

©The University of Texas at Austin, 2018

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text: "Optimization on Manycore", Texas Advanced Computing Center, 2018. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License.

