



TEXAS ADVANCED COMPUTING CENTER

WWW.TACC.UTEXAS.EDU



TEXAS

The University of Texas at Austin

Manycore Optimization: Multitasking

John Cazes

February 16, 2018

Multitasking

Basic Concepts: A task is an active program and related resources

- ▶ *Multiprocessing*: concurrently running tasks are processes
 - ▶ processes do not share memory by default
- ▶ *Multithreading*: concurrently running tasks are threads
 - ▶ threads do share memory by default
- ▶ *speedup*: runtime of serial versus parallel $\frac{T_s}{T_p}$
- ▶ *efficiency*: speedup versus number of tasks $\frac{T_s}{pT_p}$
 - ▶ Typically less than 1
- ▶ Process/memory *affinity*: pinning tasks to sets of cores and memory
- ▶ *Contention*: tasks competing for bandwidth and ownership of data
- ▶ *Dependencies*: execution of code by one task needs results from another (similar to dependency analysis in vectorization)
- ▶ *Granularity*: amount of work per task

Granularity: Load Balancing vs Overhead

Granularity is the amount of work per task

- ▶ It is important to balance granularity for parallel efficiency
- ▶ Too much granularity can lead to load imbalance
 - ▶ Some tasks could be idle for long periods, waiting for others to complete
- ▶ Too little granularity can lead to large overhead
 - ▶ Parallelizing has overhead
 - ▶ Partitioning problem
 - ▶ Task startup
 - ▶ Scheduling work to tasks
 - ▶ Communicating/sharing data between tasks
 - ▶ Lock management and synchronization

Load Imbalance Example

A loop with 10 iterations, divided among 9 threads

```
for(j=0;j<10;j++){...}
```

- ▶ Each iteration takes 1 thread 1 day
- ▶ 9 iterations are done the 1st day by 9 threads
- ▶ 1 iteration is done the 2nd day by 1 thread (9 threads are idle for a day :()
- ▶ Negligible communication overhead but poor load balance
- ▶ $Speedup = 10/2 = 5$ vs $Speedup = 10$ in perfect balance

Parallel Overhead Example

A loop with 10 iterations, divided among 10 threads

```
for(j=0; j<10; j++) { ... }
```



- ▶ Setup of parallelization requires 1ms
- ▶ Each iteration takes 1 thread 0.1ms
- ▶ 10 iterations are done in parallel in 0.1 ms
- ▶ Total parallel time is 1.1ms → serial would have run in 1ms!
- ▶ $Speedup = 1/1.1$

OpenMP

We will focus on OpenMP (OMP) here

- ▶ MPI considerations are similar on Manycore to Multicore
 - ▶ One caveat to that: Typically use less than 68 MPI tasks per node
 - ▶ MPI has a larger memory footprint than OpenMP
 - ▶ MPI composed of processes
- ▶ OpenMP is very common in scientific applications
 - ▶ OpenMP composed of threads
 - ▶ shared memory multithreading API for C,C++,Fortran
- ▶ Other approaches to parallelization exist:
 - ▶ multiprocessing, e.g. Python multiprocessing module
 - ▶ launchers
 - ▶ pthreads

OMP Overhead: EPCC Benchmark on KNL



Impact of Overhead and Imbalance on Performance

Measure in FLOPs lost

- ▶ KNL: 3.264 TF theoretical peak $\rightarrow \sim 2$ TF (e.g. HPL)
- ▶ Overhead for parallel directive: $0.58\mu s \rightarrow 11\mu s$
- ▶ $\frac{2TFLOPs}{s}(0.58 \rightarrow 11)\mu s = (1.2 \rightarrow 22)MFLOP$
- ▶ Cost of parallelization needs to be worthwhile \equiv sufficient work per task
- ▶ Short loops or small amount of work is rarely worth parallelizing
- ▶ Any reduction in overhead you can get is worth it!

Overhead reduction: Thread Pools

OMP minimizes overhead of thread creation *automatically* (GCC & Intel)

- ▶ Thread pool is created the first time a parallel construct is encountered
- ▶ Thread creation can take thousands of microseconds
- ▶ Threads from pool are reused for subsequent parallel work
- ▶ If new threads are needed they are created and added to pool
- ▶ Thread pool creation overhead occurs only in first call

```
#pragma parallel for // thread creation overhead
for(j=0;j<10;j++){...}
.
.
#pragma parallel for // no thread creation overhead
for(k=0;k<100;k+){...}
```

Loop Modifications for Parallel Performance

- ▶ Loops are a common place for parallelization
 - ▶ Divide iterations among threads
 - ▶ Iterations must be independent
- ▶ Loop modifications can make parallelization possible
 - ▶ Remove dependencies between iterations
- ▶ Loop modifications can improve granularity
 - ▶ Can improve load balancing
 - ▶ Can reduce overhead

General tip for good loop granularity in loops

Parallelize over loops with high iteration counts relative to the number of threads.

- ▶ Overhead is less significant: Time to setup versus time spent on work is small
- ▶ Load imbalance is less significant: Time for “remainder” iterations versus time all threads are working is smaller

Loop fission

Can enable parallelization by removing dependencies

Later iteration needs result from earlier iteration (RAW dependency):

```
for (j=1; j<N; j++){  
    a[j]=b[j-1];  
    b[j]+=1;  
}
```

Loop fission can remove dependency and enable parallelization:

```
#pragma parallel for  
for (j=1; j<N; j++){  
    b[j]+=1;  
}
```

```
#pragma parallel for  
for (j=1; j<N; j++){  
    a[j]=b[j-1];  
}
```

Also can improve cache reuse

Loop Fusion

Increases granularity and reduces overhead from parallelization and loop

```
#pragma parallel for
  for (j=0; j<N; j++)
    a[j]=b[j]*c[j];
#pragma parallel for
  for (j=0; j<N; j++)
    d[j]=e[j]*f[j];
```

Loop fusion

```
#pragma parallel for
  for (j=0; j<N; j++){
    a[j]=b[j]*c[j];
    d[j]=e[j]*f[j];
  }
```

Also more data is in flight which is good if BW is not saturated

Merge nested loops

Increases iteration count → better load balancing

```
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        d[i][j]=e[i]*f[j];
```

Merge

```
for (k=0; k<N*N; k++) {  
    i=k/N;  
    j=k%N;  
    d[i][j]=e[i]*f[j];  
}
```

OpenMP has a Directive for this: `collapse(#loops)`

```
#pragma omp parallel for collapse(2)  
for (i=0; i<N; i++)  
    for (j=0; j<N; j++)  
        d[i][j]=e[i]*f[j];
```

Reduce overhead

OpenMP `parallel` has some overhead (even with thread pools)

ok

```
for (rep=0; rep<M; rep++)  
  #pragma omp parallel for  
    for (j=0; j<N; j++)  
      A[j] += x[j];
```

better

```
#pragma omp parallel  
{  
  for (rep=0; rep<M; rep++)  
    #pragma omp for  
      for (j=0; j<N; j++)  
        A[j] += x[j];  
}
```

Load balancing with OpenMP Scheduling

OpenMP has 3 *scheduling* types

```
#pragma omp parallel for schedule(scheduletype, chunksize)
for(k=0;k<N*N;k++)
    work(k);
```

► static

- default scheduler
- same iterations for every thread
- very little overhead
- no load balancing

► dynamic

- assign work in chunksize (default=1)
- when thread completes, new work is assigned
- overhead from synchronization

► guided

- similar to dynamic
- initially gives large chunks of iterations to tasks
- gradually reduces chunks of iterations to min of chunksize
- sometimes overhead less than dynamic

Load balancing with OpenMP Scheduling

Scheduling impacts overhead and cache reuse

- ▶ Static has least overhead, dynamic most
 - ▶ larger chunksize can reduce overhead
- ▶ Dynamic may spoil cache reuse
 - ▶ After first iteration is performed by a thread the data may already be in the local cache for next iteration
 - ▶ Other cores may not have data in cache
- ▶ If a different thread on a different core is assigned the next iteration it may not have the data in it's local cache
- ▶ Larger chunk size can mitigate this effect
 - ▶ recall default chunksize is 1: Use with caution!

Load balancing with OpenMP `nowait`

`nowait` clause

- ▶ parallel constructs have implicit barrier at end where all threads must wait for other threads to complete
- ▶ `nowait` allows completed threads to advance and do more work
- ▶ usually need to use `dynamic` to benefit

```
#pragma omp parallel
{
  #pragma omp for nowait
  for (j=0; j<N; j++)
    A[j]+=f(x[j]);

  #pragma omp for
  for (j=0; j<N; j++)
    B[j]+=g(x[j]);
}
```

Cache Coherence, Memory Consistency, & Performance in Manycore


Shared data is a major performance limiter to parallelism

- ▶ If algorithm looks like it should scale but does not often due to this
- ▶ Data is kept consistent between cores' local caches by cache coherence and memory consistency
 - ▶ Recall KNL core has private L1 cache and shares L2 cache with one other core
 - ▶ Recall the 4 hardware threads on each core share an L1
- ▶ Consistency requires cycles and bandwidth
- ▶ Consistency can force serialization in otherwise parallel code
 - ▶ Can be explicit in program (e.g. `atomic`)
 - ▶ Can be transparent in program through *cache coherence protocol*

Parallel Contention: True Sharing

Major limitation to scalability!

- ▶ If core0 thread changes A's value, core1 thread must wait to change A
- ▶ KNL hyperthreads, e.g. 0,68,132,204, share L1 so not bad
- ▶ cores on local tile share L2 so not so bad
- ▶ cores on remote must request A over mesh



True Sharing Mitigation

Performance Improvement Factors

- ▶ Frequent writes to shared variables main bottleneck
- ▶ Frequent reads are OK
- ▶ `#pragma omp atomic` should be avoided (serializes code)
- ▶ Make one copy of value per thread if possible (even if redundant): `private` clause
- ▶ OpenMP has efficient `reduction` directive for many operators: `+`, `*`, `-`, `|` etc.

```
#pragma omp parallel for
for(i=0;i<N;i++)
    #omp atomic
    a+=i;


int tmp, result;
#pragma omp parallel private(tmp) shared(result)
{
    #pragma omp for
    for(i=0;i<N;i++) tmp+=i;
    #pragma omp critical
        result+=tmp
}
```

```
#pragma omp parallel for reduction(a,+)
for (i=0;i<N;i++)
    a+=i;
```

Parallel Contention: False Sharing

Can happen easily but sometimes hard to detect

- ▶ Cache lines hold multiple elements (64B or 8 doubles on KNL)
- ▶ Different data on same cache line may be modified by multiple threads
- ▶ Entire cache lines are invalidated by modification, not individual data elements
- ▶ Shared cache line is forwarded or written to L2/memory constantly to maintain coherency



False Sharing Mitigation

Strategies

- ▶ Allocate data used by each thread contiguously (definitely do not interleave)
- ▶ Watch out for sequentially declared variables (may be land on same cache line)
- ▶ Chunksize - make multiple of cache line
- ▶ Spreading, add a dimension to array so each thread has own cache line
 - ▶ Reduces cache reuse → increases memory bandwidth but reduces cache coherency traffic

//contiguous b.x and b.y probably on same cache line

```
struct b {  
    int x;  
    int y;  
};
```

//a access is interleaved: bad chunksize

```
double a[N];  
#pragma omp parallel for schedule(static,1)  
for (i=0;i<N;i++)  
    a[i]+=i
```

//each thread has a separate cache line: good chunksize

```
double a[N];  
#pragma omp parallel for schedule(static,8)  
for (i=0;i<N;i++)  
    a[i]+=i
```

// each thread has separate cache line: spreading

```
double b[N][8];  
#pragma omp parallel for  
for (i=0;i<N;i++)  
    b[i][0]+=i
```

Detecting Data Contention

Can be difficult: profilers are a good way to go

- ▶ If algorithm will not scale
- ▶ Looks for heavy cache traffic with using, for example `perf stat -d ./executable`
- ▶ `perf` is a light-weight tool built into newer (2.6.32+) Linux distros
- ▶ Grants access to hardware performance monitoring counters (PMCs)
- ▶ Try Intel's Vtune - also uses PMCs

perf stat -d ./executable

Example on KNL

Performance counter stats for 'numactl -C 0,2 ./sharing':

11437.839149	task-clock (msec)	#	1.822 CPUs utilized	
185	context-switches	#	0.016 K/sec	
1	cpu-migrations	#	0.000 K/sec	
408	page-faults	#	0.036 K/sec	
17,007,721,754	cycles	#	1.487 GHz	(33.34%)
<not supported>	stalled-cycles-frontend			
<not supported>	stalled-cycles-backend			
2,027,617,926	instructions	#	0.12 insns per cycle	(50.00%)
333,745,966	branches	#	29.179 M/sec	(49.99%)
590,178	branch-misses	#	0.18% of all branches	(50.01%)
<not supported>	L1-dcache-loads			
21,916,918	L1-dcache-load-misses	#	1.916 M/sec	(33.35%)
22,675,535	LLC-loads	#	1.983 M/sec	(33.33%)
<not supported>	LLC-load-misses			

6.278683864 seconds time elapsed

High BW Memory and Parallel Performance


- ▶ For BW bound MCDRAM is better
 - ▶ More threads = higher BW requirements
- ▶ For latency bound DDR can be better
- ▶ Cache mode is easy to use but Flat is more performant
 - ▶ Cache mode has extra memory latency due to caching mechanisms
 - ▶ Flat mode DDR best for latency sensitive
 - ▶ Flat mode MCDRAM best for BW sensitive
- ▶ Crossover between modes can occur as more threads are added and BW saturates

	DDR	MCDRAM
Capacity	100GB	16 GB
Bandwidth	90GB/s	450GB/s
Latency	125ns	150ns

CACHE MODE



FLAT MODE



Aside: OpenMP + Vectorization

With OpenMP 4.0: `#pragma omp simd`

- ▶ New portable directive
- ▶ Tells compiler to vectorize a loop
- ▶ more portable than intrinsics (vendor agnostic)
- ▶ can specify even a function as simd safe
- ▶ **Warning:** Inhibits compiler safety analysis!

```
#pragma omp parallel for simd
for(i=0;i<N;i++)
    a[i] += c[i]

#pragma omp declare simd
double vectorized_function(double *a, double *b){...}

#pragma omp simd
for(i=0;i<N;i++) c[i] += vectorized_function(a[i], b[i]);
```

Aside: Super-linear Speedup/Slowdown

Speed-up exceeds number of cores

- ▶ Usually happens when problem is decomposed small enough to fit in each core's cache
- ▶ Serial problem size did not fit in cache - streamed from memory and back, possibly many times
- ▶ Basically improved cache line reuse
- ▶ Be careful analyzing speed-up with datasets that are too small!

Slowdown for larger problem sizes even with more threads

- ▶ Little cache reuse
- ▶ Memory accesses are dominating
- ▶ Hide latency

License

©The University of Texas at Austin, 2018

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text: "Manycore Optimization", Texas Advanced Computing Center, 2018. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License.

