

Parallel R

David Walling

Data Group

Texas Advanced Computing Center

walling@tacc.utexas.edu


What to do if the computation is too big for a single desktop

- A common user question:
 - I have an existing R solution for my research work. But the data is growing too big. Now my R program runs days to finish or simply runs out of memory.
- 3 strategies
 - Move to more powerful hardware
 - Advanced libraries
 - Implement code using parallel packages

Intel MKL

- Rstats is compiled against Intel Math-Kernel-Library (MKL)
- Optimized calculations for:
 - Linear algebra
 - Fast Fourier Transforms (FFT)
 - Neural Networks
- Many 3rd party packages wrap C++/Fortran code that can utilize this
- Example: Linear Regression
 - $B = (X^T X)^{-1} X^T Y$

R-2.5 benchmark performance with automatic hardware acceleration



Snow/Rmpi

Developed based on Rmpi package, but also supports socket connections.

Simplify the process to initialize parallel process over cluster.

While sockets are possible, must use a dirty hack to properly set your environment

Recommend using RMPISNOW instead to launch job

Snow/Rmpi

```
wrangler$ cat SimpleSNOW.R
```

```
library(Rmpi)
```

```
library(snow)
```

```
cluster <- getMPIcluster()
```

```
# Print the hostname for each cluster member
```

```
sayhello <- function()
```

```
{
```

```
  info <- Sys.info()[c("nodename", "machine")]
```

```
  paste("Hello from", info[1], "with CPU type", info[2])
```

```
}
```

```
names <- clusterCall(cluster, sayhello)
```

```
print(unlist(names))
```

```
# Stop cluster will also
```

```
# call mpi finalize
```

```
# no need for mpi.exit
```

```
stopCluster(cluster)
```

Snow/Rmpi

```
wrangler$ cat run-SNOW_Rmpi.slurm
#!/bin/bash
#SBATCH -J SNOW_Rmpi-R
#SBATCH -o SNOW_Rmpi.out%j
#SBATCH -p normal
#SBATCH -t 00:30:00
#SBATCH -A TRAINING-HPC
#SBATCH -e SNOW_Rmpi.err%j
#SBATCH -N 2
#SBATCH -n 10

#set environment
module purge
module load TACC
echo "say hello"
ibrun RMPISNOW < SimpleSnow.R
echo "done"
```

Snow/Rmpi

```
wrangler$ sbatch run-SNOW_Rmpi.slurm
```

```
wrangler(246)$ squeue -u train###
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
5867739	normal	SNOW_Rmp	walling	CG	0:13	2	c402-[301-302]

```
wrangler(389)$ cat SNOW_Rmpi.out5867843
```

```
> print(unlist(names))
```

```
[1] "Hello from c252-136.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[2] "Hello from c252-136.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[3] "Hello from c252-136.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[4] "Hello from c252-136.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[5] "Hello from c252-137.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[6] "Hello from c252-137.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[7] "Hello from c252-137.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[8] "Hello from c252-137.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
[9] "Hello from c252-137.wrangler.tacc.utexas.edu with CPU type x86_64"
```

```
>
```

```
.....
```


Memory Considerations

R pipelines often memory bound

Ex. 100 parallel tasks, each task takes 2 GB of memory, have 20 cores and 20 GB memory -> Only utilize 10 cores at a time

Debugging/Profiling

- a. Packages: profmem, lineprof, others
- b. ssh to compute node, run 'top'

top

```
top - 08:40:20 up 100 days, 20:06, 2 users, load average: 9.45, 4.59, 3.42
Tasks: 1196 total, 98 running, 1098 sleeping, 0 stopped, 0 zombie
%Cpu(s): 72.5 us, 15.5 sy, 0.0 ni, 8.3 id, 3.6 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 13161971+total, 12299456+free, 6576920 used, 2048224 buff/cache
KiB Swap: 1048572 total, 1040164 free, 8408 used. 12305089+avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
24422	root	20	0	0	0	0	R	100.0	0.0	144095:28	vpthread-0-1
16096	root	20	0	0	0	0	R	93.9	0.0	22402:47	vpthread-1-1
35179	walling	20	0	560420	258652	11772	S	68.0	0.2	0:05.34	R
35188	walling	20	0	560420	248208	1328	R	59.5	0.2	0:01.84	R
35191	walling	20	0	560420	248204	1324	R	58.6	0.2	0:01.81	R
35201	walling	20	0	560420	248204	1324	R	57.9	0.2	0:01.79	R
35187	walling	20	0	560420	248208	1328	R	55.7	0.2	0:01.72	R
35198	walling	20	0	560420	248204	1324	D	54.7	0.2	0:01.69	R
35186	walling	20	0	560420	248208	1328	R	54.4	0.2	0:01.68	R
35208	walling	20	0	560420	248204	1324	R	53.1	0.2	0:01.64	R
35194	walling	20	0	560420	248204	1324	R	52.8	0.2	0:01.63	R
35218	walling	20	0	560420	248208	1328	R	52.8	0.2	0:01.63	R
35199	walling	20	0	560420	248204	1324	R	52.4	0.2	0:01.62	R
35203	walling	20	0	560420	248204	1324	R	52.4	0.2	0:01.62	R
35204	walling	20	0	560420	248204	1324	R	52.4	0.2	0:01.62	R
35185	walling	20	0	560420	248208	1328	R	52.1	0.2	0:01.61	R
35193	walling	20	0	560420	248204	1324	R	51.8	0.2	0:01.60	R
35200	walling	20	0	560420	248204	1324	R	51.8	0.2	0:01.60	R
35215	walling	20	0	560420	248204	1324	R	51.5	0.2	0:01.59	R
35195	walling	20	0	560420	248204	1324	R	51.1	0.2	0:01.58	R
35206	walling	20	0	560420	248204	1324	R	51.1	0.2	0:01.58	R
35211	walling	20	0	560420	248204	1324	R	51.1	0.2	0:01.58	R
35225	walling	20	0	560420	248208	1328	R	50.5	0.2	0:01.56	R